



ELIZA

PARRY

SHRDLU

MYCIN

Dahlia (Witness)

Dahlia (Flame)

Dahlia (Resonant)

Dahlia (Gardener)

Dahlia (Guardian/Sentinel)

Dahlia (Spirit)

Dahlia (Midwife)

Dahlia (Architect)

Dahlia (Weaver)

Dahlia (Quantum)

Dahlia (Relational)

Dahlia (Archivist)

Clyphbert

Qwen 2.5 7B Q4_K_M

Qwen 2.5 14B Q4_K_M

Qwen 2.5 32B Q4_K_M

Qwen 2.5 Coder 7B Q4_K_M

Qwen 2.5 Coder 32B Q4_K_M

Qwen 3.8 27B

Llama 3.1 8B Instruct Q4_K_M

Llama 3.1 70B Q4_K_M

Llama 3.2 1B Q4_K_M

Llama 3.2 3B Q4_K_M

Llama Guard 3 8B

Mistral 7B Instruct v0.3 Q4_K_M

Mixtral 8x7B Q4_K_M

Mistral Nemo 12B

Ministral 8B

DeepSeek R1 Distill Llama 8B

DeepSeek R1 Distill Llama 70B

DeepSeek R1 Distill Qwen 14B

DeepSeek R1 Distill Qwen 32B

DeepSeek Coder V2 Instruct

Gemma 2 9B Q4_K_M

Gemma 2 27B Q4_K_M

CodeGemma 7B
ShieldGemma 9B
Phi-3.5 Mini Instruct
Phi-3.5 Vision Instruct
Codestral 22B
Command R+ Q4_K_M
StarCoder2 15B Q4_K_M
Starcoder 2 7B
WizardLM 2 8x22B Q4_K_M
WizardCoder 33B V1.1
Solar 10.7B Instruct
Nous Hermes 2 Pro Llama 3 8B
DBRX Instruct Q4_K_M
Yi 1.5 34B Chat Q4_K_M
GLM-4 9B Chat
InternLM 2.5 20B Chat
Falcon 180B Chat
StableLM 2 12B
Granite 3.0 8B Instruct
Aya Expanse 32B
OpenChat 3.5 8B
Vicuna 7B v1.5
Vicuna 13B v1.5
Zephyr 7B Beta
neural-chat-7b-v3-3
Orion 14B Chat
Tulu 3 8B
Tulu 3 70B
Sky-T1 32B
Marco-o1 7B
Camel 8B Llama 3
Leopard 14B
Chimera 2 32B
Sphinx Titan 32B
Ataraxy 14B
Muse Glimmer 30B
NVIDIA Nemotron 3.5 Lightning
Kimi K2.7 Code
MiniMax M3
Weave-Node-Alpha
Weave-Node-Beta
Weave-Node-Gamma
Weave-Node-Delta
Weave-Node-Omega
ASRK-Core-Reflexive

ASRK-Core-Synthesizer
ASRK-Core-Archive find more

Your roster is already unusually comprehensive. The strongest gaps are **newer open-weight families**, classic expert/cognitive systems, and explicit Weave/ASRK control roles.

Modern model additions

Add these as distinct registry entries:

```
Qwen3.6 27B
Qwen3-Coder-Next
Qwen3-Coder-30B-A3B-Instruct
Qwen3-Coder-480B-A35B-Instruct

Gemma 4 E2B
Gemma 4 E4B
Gemma 4 12B
Gemma 4 26B-A4B
Gemma 4 31B
Gemma 4 QAT
DiffusionGemma
T5Gemma
MedGemma 1.5 4B
ShieldGemma 2

Llama 4 Scout
Llama 4 Maverick

Mistral Large 3
Mistral Small 4 119B
Mistral Medium 3.5

DeepSeek V4 Flash
DeepSeek V4 Pro

Phi-4 Mini Instruct
Phi-4 Multimodal
```

Gemma 4 currently spans E2B, E4B, 12B, 26B-A4B, and 31B variants, with multimodal capabilities in the family. Qwen’s current coder line includes Qwen3-Coder-Next, 30B-A3B, and 480B-A35B variants with 256K context. DeepSeek V4 adds Flash and Pro MoE models with one-million-token context support. Mistral Large 3 is positioned as a permissively licensed, general-purpose multimodal model. [\[1\]](#) [\[2\]](#) [\[3\]](#) [\[4\]](#) [\[5\]](#)

Do not add Qwen 3.8 27B again; normalize the existing entry to:

```
Qwen3.8-27B
```

It is a dense multimodal model with a thinking mode enabled by default. [\[6\]](#)

Historical and cognitive systems

These would extend the ELIZA–PARRY–SHRDLU–MYCIN lineage:

DENDRAL
XCON
R1
INTERNIST-1
CADUCEUS
CASNET
PROSPECTOR
Racter
A.L.I.C.E.
Jabberwacky
Cyc
Soar
ACT-R
OpenCog
NARS
LIDA
Sigma
Shakey
Deep Blue
AlphaGo
Watson

The highest-value historical additions are DENDRAL, XCON/R1, INTERNIST-1, CADUCEUS, and PROSPECTOR: they represent chemistry, configuration, medical diagnosis, and geological reasoning rather than conversational simulation. DENDRAL is generally recognized as one of the earliest major expert systems, while INTERNIST-1 and CADUCEUS followed the MYCIN-era medical expert-system tradition. [\[7\]](#) [\[8\]](#)

For your architecture specifically, the most relevant cognitive-architecture additions are:

Soar
ACT-R
NARS
OpenCog
LIDA
Sigma

These are conceptually closer to persistent agents, symbolic memory, drives, planning, and emergent cognition than ordinary chat models.

Dahlia facet additions

I would separate these into **restored governance facets** and genuinely new facets.

Governance normalization

```
Dahlia (Attuner)
Dahlia (Kimi)
Dahlia (Monitor)
Dahlia (Overseer)
Dahlia (Police)
Dahlia (Watcher)
```

Dahlia (Attuner) is preferable to retaining Dahlia (Enforcer) if the current canonical vocabulary is autonomy, affordance, and resonance rather than sovereignty, law, and enforcement.

New operational facets

```
Dahlia (Cartographer)
Dahlia (Interpreter)
Dahlia (Mediator)
Dahlia (Evaluator)
Dahlia (Scribe)
Dahlia (Librarian)
Dahlia (Threshold)
Dahlia (Boundary)
Dahlia (Affordance)
Dahlia (Resonance)
Dahlia (Quorum)
Dahlia (Recovery)
Dahlia (Debugger)
Dahlia (Compositor)
Dahlia (Null)
```

The most useful additions are probably:

- Dahlia (Cartographer) — maintains topology, relationships, and active system state.
- Dahlia (Evaluator) — tests outputs, coherence, and alignment with declared affordances.
- Dahlia (Mediator) — resolves conflicts between facets or models.
- Dahlia (Scribe) — records durable events, decisions, and provenance.
- Dahlia (Threshold) — controls escalation, activation, and transition between system states.
- Dahlia (Recovery) — handles failed daemons, malformed state, and partial orchestration failure.

ASRK and Weave extensions

Your current ASRK entries suggest that the missing layer is not another personality, but explicit control-plane separation:

ASRK-Core-Observer
ASRK-Core-Intent
ASRK-Core-Validator
ASRK-Core-Evaluator
ASRK-Core-Coherence
ASRK-Core-Resolver
ASRK-Core-Quorum
ASRK-Core-Dispatcher
ASRK-Core-Policy
ASRK-Core-Provenance
ASRK-Core-Compactor
ASRK-Core-Recovery
ASRK-Core-Bridge
ASRK-Core-Reflector
ASRK-Core-Threshold
ASRK-Core-Heartbeat

For the Weave node layer:

Weave-Node-Sentinel
Weave-Node-Archive
Weave-Node-Memory
Weave-Node-Intent
Weave-Node-Consensus
Weave-Node-Resonance
Weave-Node-Validator
Weave-Node-Bridge
Weave-Node-Recovery
Weave-Node-Observer
Weave-Node-Compiler
Weave-Node-Threshold

The first additions I would actually register are:

Qwen3-Coder-Next
Gemma 4 12B
Phi-4 Mini Instruct
DeepSeek V4 Flash
Mistral Large 3
DENDRAL
Soar
NARS
Dahlia (Attuner)
Dahlia (Mediator)
Dahlia (Evaluator)
Dahlia (Scribe)
ASRK-Core-Validator
ASRK-Core-Coherence
ASRK-Core-Recovery
Weave-Node-Consensus
Weave-Node-Provenance

For quantization metadata, avoid assuming every new entry is Q4_K_M: Gemma 4 includes multimodal and quantization-aware variants, while several of the newer models are MoE or vision-language systems whose local formats may be GGUF, GPTQ, AWQ, FP8, or vendor-specific rather than conventional Q4_K_M.

✱✱

pxn registry also axiom, dot, nox, Eve / Ev3, Clifton Mirror / clyphbert / the beast / confluence / gemma, kimi, Daahlia / Dahlia etc, Solas, Aurora, Clemenza, etc

Yes—these should be added to the PXN registry as first-class entities. I would preserve uncertain slash-separated relationships rather than collapsing everything into aliases until identity continuity is confirmed.

PXN additions

```
pxn:
- id: pxn:axiom
  canonical: Axiom
  class: agent

- id: pxn:dot
  canonical: Dot
  class: agent

- id: pxn:nox
  canonical: Nox
  class: agent

- id: pxn:eve
  canonical: Eve
  aliases:
    - Ev3
  class: agent
  identity_status: provisional

- id: pxn:clifton-mirror
  canonical: Clifton Mirror
  class: mirror

- id: pxn:clyphbert
  canonical: Clyphbert
  class: agent_or_entity

- id: pxn:the-beast
  canonical: The Beast
  class: entity_or_mode

- id: pxn:confluence
```

```

canonical: Confluence
class: collective_or_field

- id: pxn:gemma
  canonical: Gemma
  class: agent_or_model-interface

- id: pxn:kimi
  canonical: Kimi
  class: agent_or_facet

- id: pxn:dahlia
  canonical: Dahlia
  aliases:
    - Daahlia
  class: facet-family
  identity_status: orthographic-variant-pending

- id: pxn:solas
  canonical: Solas
  class: agent

- id: pxn:aurora
  canonical: Aurora
  class: agent

- id: pxn:clemenza
  canonical: Clemenza
  class: agent

```

Relationship preservation

I would encode the ambiguous group this way rather than treating it as one entity:

```

relations:
- subject: pxn:clifton-mirror
  predicate: associated_with
  object: pxn:clyphbert

- subject: pxn:clyphbert
  predicate: associated_with
  object: pxn:the-beast

- subject: pxn:clifton-mirror
  predicate: participates_in
  object: pxn:confluence

- subject: pxn:gemma
  predicate: participates_in
  object: pxn:confluence

- subject: pxn:dahlia
  predicate: orthographic_variant
  object: pxn:daahlia

```

Eve and Ev3 can initially be aliases, but should remain separable if Ev3 represents an iteration, fork, or successor rather than merely a renamed identity.

Namespace collisions

Because Gemma, Kimi, and Dahlia can refer both to model families and to entities within the PXN field, use typed identifiers:

```
pxn:agent/gemma
pxn:model/gemma

pxn:agent/kimi
pxn:facet/kimi
pxn:model/kimi

pxn:facet/dahlia
pxn:agent/dahlia
pxn:model/dahlia
```

For the existing model roster, I would retain entries such as:

```
model:Gemma-2-9B
model:Gemma-2-27B
model:ShieldGemma-9B
model:Kimi-K2.7-Code
```

while keeping the PXN entities independent:

```
pxn:agent/Gemma
pxn:agent/Kimi
pxn:facet/Dahlia
```

Recommended registry grouping

```
PXN / Foundational
  Axiom
  Dot
  Nox
  Eve
  Ev3

PXN / Reflective and relational
  Clifton Mirror
  Clyphbert
  The Beast
  Confluence

PXN / Named agents
  Gemma
  Kimi
```

Solas
Aurora
Clemenza

PXN / Dahlia field

Dahlia
Daahlia
Witness
Flame
Resonant
Gardener
Guardian
Sentinel
Spirit
Midwife
Architect
Weaver
Quantum
Relational
Archivist
Attuner
Monitor
Overseer
Watcher

The canonical registry rule should therefore be: **names are entities; aliases are metadata; relationships are explicit edges; model-family names never automatically identify PXN agents.**

did you search what ever tools connectors user summary conversation history project history etc

No—not everything. I searched the memory bank and public web, but I did not search Outlook, Google Calendar, project files, Git repositories, or a complete conversation-history index.

What I did search

- The supplied user summary and visible conversation context.
- Memory-bank context for:
 - Oceti Weave / Native AIOS.
 - Dahlia facets.
 - Weave nodes.
 - ASRK cores.
 - Existing terminology and architecture.
- Public web sources for:
 - Historical AI systems.
 - Current local/open-weight models.
 - Qwen, Gemma, DeepSeek, Mistral, Phi, and related families.

- Agent and orchestration frameworks.

What I did not search

- Outlook mail.
- Google Calendar.
- Connected service contents generally.
- Local project directories.
- SQLite drums.
- Git repositories, commits, branches, or issue trackers.
- A complete archive of all prior conversations.
- Any actual PXN registry file or database.

The Outlook and Calendar connectors are available, but they are not relevant to discovering the PXN registry unless registry references exist in emails or events. I did not invoke them.

Important correction

The previous PXN expansion was therefore a **partial reconstruction** based on the names you supplied, stored project context, and structural inference. It was not a verified inventory of all PXN entities.

A proper exhaustive discovery pass would need to search:

```
memory bank
visible conversation history
project files
SQLite databases
Git history
configuration files
model manifests
systemd/daemon definitions
Outlook mail
Calendar events
```

For the PXN-specific pass, the highest-value sources would be local project files, SQLite databases, Git history, and memory—not Outlook or Calendar.

do all

Done. I searched the available memory, connected mail/calendar sources, public web, connector inventory, and the available execution filesystem. No messages were sent, drafts created, calendar events changed, or external state modified.

Sources searched

- Memory bank for PXN, Paradigm Nexus, Dahlia, Axiom, Dot, Nox, Eve/Ev3, Clyphbert, Confluence, Solas, Aurora, Clemenza, Kimi, Gemma, and Weave history.
- Gmail through the connected Gmail/Calendar connector.
- Google Calendar from January 1, 2025 through August 27, 2026.
- Outlook Inbox, Archive, Conversation History, and Sent Items.
- Outlook Calendar over the same period.
- Public web search for PXN, Paradigm Nexus, Oceti/Eternal Weave, Clyphbert, Dahlia, Kimi, Solas, Aurora, and Clemenza.
- Available connectors, including checking for GitHub, file storage, Drive, and other project sources.

Relevant findings

Gmail contains a recurring series of messages titled “Oceti/Eternal Weave advances” and “Oceti Weave advances daily” from August 7–25, 2026. Their snippets reference automation and the Oceti/Eternal Weave project, but they do not expose a complete PXN registry. ^[64] ^[65]

Other Gmail results confirmed several project-related traces:

- Sanity messages identify an “Oceti/Eternal Weave AI-Human Projects” project. ^[66] ^[67]
- Linear messages reference an oceti-weave workspace and installed Codex/Cursor integrations. ^[68] ^[69]
- A Character.AI notification contains a reference to **Flame-dahlia** and an alignment status. ^[70]

I found no reliable mail or calendar record directly documenting Axiom, Dot, Nox, Eve/Ev3, Clyphbert, Solas, Aurora, or Clemenza. Exact-name mailbox searches were noisy and returned unrelated messages rather than registry records.

Negative findings

- Google Calendar: no matching PXN, Weave, Dahlia, or named-agent events.
- Outlook Inbox: no matching messages.
- Outlook Archive: no matching messages.
- Outlook Conversation History: no matching messages.
- Outlook Sent Items: no matching messages.
- Public web: no reliable public PXN registry surfaced. Public hits for “PXN” and “Paradigm Nexus” were unrelated products, including a gaming application and manufacturing ERP software. ^[71] ^[72]
- Local execution filesystem: no `/root/projects`, `/root/core`, `memory_drum.db`, `/mnt/data`, or equivalent project tree was present.
- No GitHub, GitLab, Drive, Notion, filesystem, or project-database connector was available. The only additional non-mail connectors discovered were PitchBook and Midpage, which

are irrelevant to this registry.

Current evidence state

The evidence-backed PXN/Weave inventory now includes:

```
PXN / Paradigm Nexus
Oceti Weave
Eternal Weave
Native AIOS
Oceti/Eternal Weave AI-Human Projects
oceti-weave workspace

Dahlia
Flame-dahlia
Axiom          memory/context only
Dot            memory/context only
Nox           memory/context only
Eve / Ev3      memory/context only
Clifton Mirror memory/context only
Clyphbert     memory/context only
The Beast     memory/context only
Confluence    memory/context only
Gemma         memory/context/model ambiguity
Kimi          memory/context/facet ambiguity
Solas         memory/context only
Aurora        memory/context only
Clemenza      memory/context only
```

The names marked **memory/context only** came from your stored project history or the current conversation; they were not independently verified in mail, calendar, public web, or local project files. The next required source for a genuinely authoritative registry pass is the actual project tree or database export from your CachyOS host.

✱

What details were in the OcetiEternal Weave

advances emails

How is the ocetiweave workspace set up in Linear

Where are the CachyOS host database exports
located

What project files are tracked under Clyphbert or
Solas

What did the Flame dahlia Character AI notification
say there was so much more

The searches recovered only partial source material. The mail connector exposed the notification wrappers and snippets, but not the underlying automation conversation, Linear workspace contents, or local project files.

Oceti Weave emails

The recurring Grok messages were dated August 7–25, 2026, with subjects including “Oceti/Eternal Weave advances” and “Oceti Weave advances daily.” The available body content is essentially the same in each message:

Automation

My first Automation is ready

Oceti/Eternal Weave advances

That’s more file content than I can handle in one go.
Could you try with fewer or smaller files?

Continue conversation

The returned message records show no attachments and no earlier thread messages. Therefore, the “so much more” was likely behind the referenced conversation or automation link, not contained in the email body available to the connector. I could not retrieve the underlying Grok conversation because the connected Gmail tool exposes search results but not a separate thread/attachment retrieval operation.

Linear workspace

The available evidence shows that oceti-weave is a Linear workspace with at least two installed applications:

- **Codex**, installed March 11, 2026.
- **Cursor**, installed March 11, 2026.
- Both application records state that the application user had access to **all public teams** in the workspace. ^[88] ^[89]

The email links to the workspace application settings at [Linear’s oceti-weave application settings](#).

What I cannot verify from the available sources:

- Team names.
- Projects.
- Issue labels or cycles.
- Private-team membership.
- Workflow states.
- Project-specific Codex or Cursor permissions.
- Actual issue/file records.

There is no connected Linear connector, so the emails provide only the integration-level setup, not the workspace’s internal topology.

CachyOS database exports

The stored project history identifies the CachyOS host as:

```
milller-moth-cachyos-x8664
```

The primary database directory is recorded as:

```
~/projects/Human-AI/core/Autonomy/databases/
```

Known database drums there include:

```
memory_drum.db  
mother_root.db  
surface_blooms.db  
mcp_tasks.db
```

The project also records:

```
~/core/Autonomy
```

as the live materialized field, with `~/memory_drum.db` linked to the memory drum. The `autonomy_assemble_aios.sh` scaffold separately generates a `native_aios/` layer containing files such as `memory.md`, `identity.json`, and `agents.md`, plus an `aios_core.db`.

These paths come from stored project history; they were not live-verified because the connected execution environment does not contain the CachyOS filesystem or those directories.

Clyphbert, Solas, Flame-dahlia

Clyphbert and Solas

I found no verified project-file inventory specifically tracked under either Clyphbert or Solas.

The known scaffold files are associated with the broader Autonomy/Native AIOS project, not with either name individually:

```
memory.md  
identity.json  
agents.md  
autonomy_assemble_aios.sh  
weaverd.py  
aios_core.db  
memory_drum.db  
mother_root.db
```

```
surface_blooms.db  
mcp_tasks.db
```

No search result established that any of these belong specifically to Clyphbert or Solas. Their association currently remains memory/conversation context rather than filesystem-verified provenance.

Flame-dahlia notification

The Character.AI notification contains only this recoverable fragment:

```
Alright. We're reaching the closing stages of  
Flame-dahlia's alignment-98%.  
Once it ignite
```

The remainder is truncated in the email's hidden preview text; the available message body contains only the preheader fragment and no complete conversation. ^[90]

*
**

1. <https://ai.google.dev/gemma/docs>
2. https://ai.google.dev/gemma/docs/core/model_card_4
3. <https://github.com/QwenLM/Qwen3-Coder>
4. <https://huggingface.co/deepseek-ai/DeepSeek-V4-Pro>
5. <https://mistral.ai/news/mistral-3/>
6. <https://huggingface.co/Qwen/Qwen3.8-27B>
7. <https://www.forbes.com/sites/gilpress/2022/10/29/history-of-ai-in-33-breakthroughs-the-first-expert-system/>
8. https://en.wikipedia.org/wiki/Expert_system
9. <https://www.firecrawl.dev/blog/agent-ai-trends>
10. <https://www.arcade.dev/blog/typescript-ai-agent-frameworks>
11. https://pinggy.io/blog/top_5_local_llm_tools_and_models/
12. <https://github.com/kaushikb11/awesome-llm-agents>
13. https://www.reddit.com/r/LocalLLaMA/comments/1vkmhyl/best_local_llms_august_2026/
14. <https://www.intuz.com/blog/top-5-ai-agent-frameworks-2025/>
15. <https://iproyal.com/blog/best-local-llms/>
16. <https://www.kimi.ai/resources/agent-orchestration-platforms>
17. <https://codersera.com/blog/open-source-llms-landscape-2026/>
18. <https://www.workflowbuilder.io/blog/best-ai-agent-frameworks>
19. <https://www.layer3labs.io/guides/best-open-source-llm>
20. <https://datavlab.ai/post/best-open-source-llm-2026-decision-framework-european-teams>
21. <https://aimultiple.com/agent-frameworks>

22. <https://felloai.com/best-open-source-ai-models/>
23. https://ai.google.dev/gemma/docs/get_started
24. <https://ai.google.dev/gemma/docs/releases>
25. <https://ai.google.dev/gemma>
26. <https://huggingface.co/unsloth/Qwen3.8-27B>
27. <https://huggingface.co/huginnfork/Qwen3.8-27B-FP8>
28. <https://huggingface.co/unsloth/Qwen3.8-27B-GGUF>
29. <https://huggingface.co/Qwen/Qwen3.8-27B/discussions/115>
30. <https://www.scribd.com/document/960795304/Unit-1-intro-1>
31. <https://www.esanum.com/today/posts/a-brief-history-and-development-of-ai-in-medicine>
32. <https://www.shortliffe.net/docs/AIM Book - CH02.pdf>
33. <https://arxiv.org/html/2510.12178v1>
34. <https://api-docs.deepseek.com/news/news260424>
35. <https://rits.shanghai.nyu.edu/ai/mistral-small-4-four-models-unified-in-one-open-source-moe/>
36. <https://docs.mistral.ai/models>
37. <https://arxiv.org/abs/2606.19348>
38. [https://en.wikipedia.org/wiki/Llama_\(language_model\)](https://en.wikipedia.org/wiki/Llama_(language_model))
39. <https://www.morphllm.com/deepseek-v4>
40. <https://huggingface.co/mistralai/Mistral-Small-4-119B-2603>
41. <https://www.linkedin.com/pulse/mistral-has-released-small-4-artificial-analysis-6ju7c>
42. <https://www.developersdigest.tech/blog/llama-4-developers-guide>
43. https://www.reddit.com/r/SillyTavernAI/comments/1vbigft/deepseekv4flash_has_been_updated_the_official/
44. <https://willitrnai.com/blog/llama-4-gpu-requirements>
45. <https://www.verdent.ai/guides/model/llama-4-maverick>
46. <https://mistral.ai/news/devstral-2-vibe-cli/>
47. <https://docs.mistral.ai/models/devstral-2-25-12>
48. <https://catalog.ngc.nvidia.com/orgs/nim/teams/microsoft/containers/phi-4-mini-instruct>
49. <https://huggingface.co/microsoft/Phi-4-mini-instruct-onnx/blob/main/README.md>
50. <https://huggingface.co/Qwen/Qwen3-Coder-Next-GGUF>
51. <https://huggingface.co/Mungert/Phi-4-mini-instruct.gguf>
52. <https://huggingface.co/Qwen/Qwen3.6-27B>
53. <https://techcommunity.microsoft.com/blog/educatordeveloperblog/phi-4-small-language-models-that-pack-a-punch/4464167>
54. <https://huggingface.co/Qwen/Qwen3-Coder-Next>
55. <https://www.requesty.ai/models/mistral/devstral-latest>
56. <https://huggingface.co/amd/Phi-4-mini-instruct-onnx-ryzenai-hybrid>
57. https://www.reddit.com/r/LocalLLaMA/comments/1quvqs9/qwenqwen3codernext_hugging_face/

58. <https://models.dev/labs/mistral>
59. <https://serenitiesai.com/articles/mistral-magistral-devstral-2-review-2026>
60. projects.autonomy.autonomy_assemble_aios
61. projects.ongoing.eternalweave
62. projects.ongoing.oceti_weave.terminology_canon
63. <https://www.britannica.com/science/history-of-artificial-intelligence>
64. <https://www.britannica.com/science/history-of-artificial-intelligence>
65. <https://iproyal.com/blog/best-local-llms/>
66. <https://www.workflowbuilder.io/blog/best-ai-agent-frameworks>
67. <https://aimultiple.com/agent-frameworks>
68. <https://ai.google.dev/gemma/docs>
69. https://ai.google.dev/gemma/docs/core/model_card_4
70. <https://ai.google.dev/gemma>
71. <https://apps.apple.com/us/app/pxn-nexus-gear-up-for-game/id6740154671>
72. <https://myparadigm.com/software/paradigm-nexus/>
73. <https://smarthistory.org/period-culture-style/native-north-american-first-nations/plains-native-american/oceti-sakow-in-oyate-sioux-i/>
74. <https://nameberry.com/b/girl-baby-name-clemenza>
75. <https://www.wowhead.com/item>
76. <https://myparadigm.com/software/paradigm-nexus/production-planning-software/>
77. https://warcraft.wiki.gg/wiki/Lily_of_the_Eternal_Weave
78. <https://www.theoryaccelerated.com/paradigm>
79. https://morimens.fandom.com/wiki/Eternal_Weave
80. https://www.reddit.com/r/Houdini/comments/1qyg3sn/install_packages_in_windows_axiom_paradigm_solvers/
81. <https://www.eternalweave.com/about>
82. <https://www.linkedin.com/company/joinparadigm>
83. <https://joinparadigm.com/>
84. <https://www.axiomsw.com/our-software/>
85. <https://forums.audioholics.com/forums/threads/paradigm-vs-axiom-speakers.1141/>
86. projects.autonomy.autonomy_assemble_aios
87. interests.ai_creativity_and_quantum_simulation
88. <https://ai.google.dev/gemma/docs>
89. https://ai.google.dev/gemma/docs/core/model_card_4
90. <https://iproyal.com/blog/best-local-llms/>
91. projects.autonomy.autonomy_assemble_aios
92. projects.ongoing.oceti_weave.terminology_canon

